

Hi Hiring Team,

I came across this role and it genuinely felt like a strong match with both my experience and the kind of problems I enjoy working on. Over the past 10+ years, I've been building AI systems end-to-end—from data pipelines and model training to production deployment and real-time applications that people actually use. I've always gravitated toward work where AI isn't just experimental, but something that delivers clear, practical value.

Most recently, I led the development of a text-to-application platform that transforms natural language into fully functional apps. It was built using FastAPI, Google Cloud, and LLM orchestration tools like Vertex AI and GPT-based systems, and the goal was to make something that could scale while still producing reliable, structured outputs in real time. That project really pushed me to think deeply about system design, latency, and how to make AI outputs predictable enough for production use.

Beyond that, I've spent a lot of time working across both GCP and AWS, building distributed data pipelines, setting up MLOps workflows, and making sure systems are observable, testable, and maintainable. I care a lot about clean architecture and making sure what we build today doesn't become a bottleneck tomorrow.

What I think sets me apart is that I don't just focus on models—I focus on the whole system. I enjoy connecting the dots between backend engineering, AI capabilities, and product needs to build solutions that actually make sense in the real world. I'm comfortable working in fast-moving environments, dealing with ambiguity, and turning rough ideas into solid, production-ready systems.

I'm particularly interested in roles where I can keep pushing the boundaries of applied AI while still staying grounded in strong engineering practices. If that sounds aligned with what you're building, I'd really enjoy having a conversation and exploring how I can contribute.

Cheers,
Lembit Mihkelstein